

Mansoor Nabawi, #309498

Tutorial 2, 05.13.2022

Exercise 1: Data cleaning and text tokenization (5 points)

Report 1

Flowchart of the processes

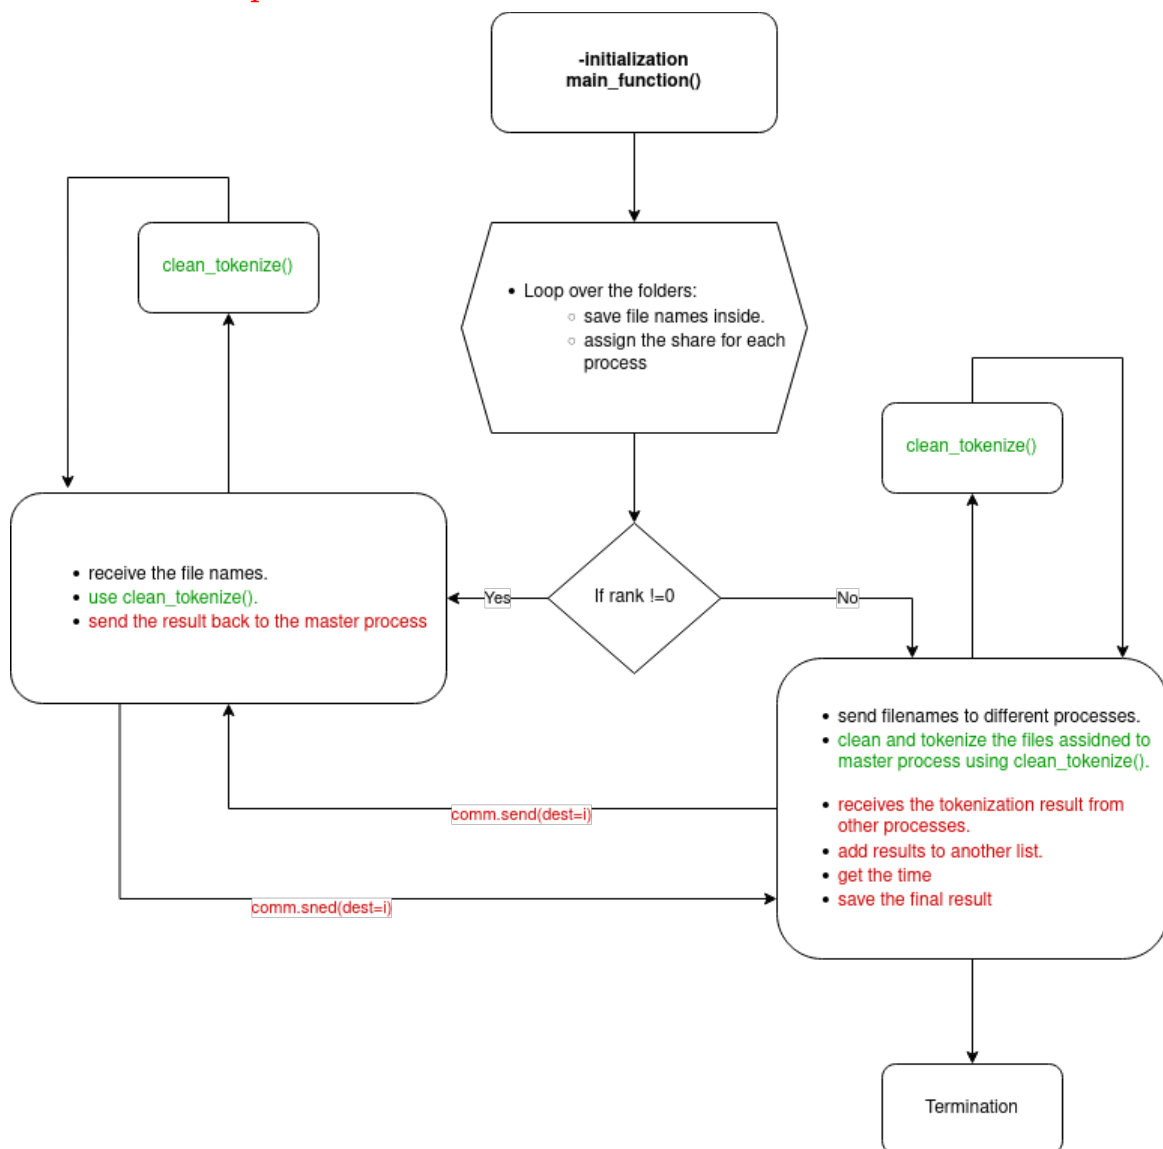


Table of Times for 1

P=1	P=2	P=3	P=4	P=6	P=8	P=10
80.93	43.68	32.3	25.73	39.0	35.5	40

- The result of parallelization is very obvious in this exercise.
- As the number of workers increases, the computation time decreases until $p=4$.

- our best result is when we use 4 workers.

Strategy

- We loop over the files of each folder.
- we share the files in each folder to different workers.
- each worker will perform cleaning and tokenization on the subset of data it has and send the result back to rank0.
- rank0 is our master process and it also work as worker, but at the end it receives all the data and add them in one big list.

result of exercise 1

```
[[['path', 'cantaloupe', 'news', 'harvard', 'near', 'howland', 'reston', 'james', 'felder',
'newsgroups', 'atheism', 'subject', 'help', 'god', 'court', 'date', 'apr', 'gmt', 'organiza
tion', 'lewis', 'resaerch', 'center', 'lines', 'distribution', 'world', 'message', 'id', 'r
eferences', 'tan', 'reply', 'nntp', 'posting', 'host', 'hopper', 'article', 'tan', 'andrew
', 'newell', 'tan', 'writes', 'article', 'apr', 'daffy', 'mccullou', 'snake', 'mark', 'mccu
llough', 'says', 'article', 'monack', 'helium', 'monack', 'helium', 'gas', 'uug', 'arizona
', 'david', 'monack', 'writes', 'another', 'issue', 'request', 'required', 'recite', 'help
', 'god', 'part', 'oath', 'theistic', 'jury', 'may', 'prejudiced', 'testimony', 'even', 'th
ough', 'atheism', 'probably', 'relevant', 'case', 'recommended', 'procedure', 'requesting',
'alternate', 'oath', 'affirmation', 'dave', 'sorry', 'using', 'follow', 'respond', 'server
', 'dropped', 'weeks', 'worth', 'news', 'keep', 'asked', 'swear', 'help', 'god', 'say', 'as
k', 'one', 'jesus', 'allah', 'vishnu', 'zues', 'odin', 'get', 'specific', 'obnoxious', 'hum
bly', 'ask', 'quitely', 'sit', 'back', 'watch', 'fun', 'james', 'felder', 'sverdrup', 'tech
nology', 'inc', 'phone', 'lewis', 'research', 'center', 'cleveland', 'ohio', 'email', 'jfel
der', 'people', 'drink', 'fountain', 'knowledge', 'people', 'gargle'], ['xref', 'cantaloupe
', 'atheism', 'talk', 'religion', 'misc', 'talk', 'origins', 'newsgroups', 'atheism', 'talk
', 'religion', 'misc', 'talk', 'origins', 'path', 'cantaloupe', 'news', 'harvard', 'near',
```

In [11]:

```

%%writefile DDA02_ex1.py

#importing libraries
import os
import numpy as np
import re
import pickle
#import nltk
from mpi4py import MPI

#loading manual stopwords
from stop_words_man import stopw
stopword_manually = stopw() #created stopwords manually

#initialization
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

def clean_tokenized(file_n):
    """
    This function receives a list of file names and goes through each file
    unnecessary words and symbols, it also checks every single word not to
    at the end tokenize them by tputting them in a list.
    args:
        - file_n: file names
    output:
        - tokenized words (list)
    """
    #empty lists will be used for final addtion and also filtering words.
    tokenized_words = []
    filtered_line = []
    for file in (file_n):
        with open(file, 'r', encoding = 'Latin-1') as f:
            text = f.read()
            #not word replaced by space
            text = [re.sub("[^a-zA-Z]", " ", f) for f in text.split()]
            text = " ".join(text)#to remove white spaces
            text = text.lower()
            text = [f for f in text.split() if len(f)>1] #single words deleted
            filtered_line = [w for w in text if not w in stopword_manually]
            tokenized_words += [filtered_line] #tokenization

    #with open("/home/mansoor/Desktop/DDA/ex02/saved"+f'/tokenized_rank_{rank}') as fp:
    #    pickle.dump(tokenized_words, fp)

    return tokenized_words

def main_function():
    """
    This function is the main function to run the parallel program on.
    """
    #time
    t0 = MPI.Wtime()

    path = "/home/mansoor/Desktop/DDA/ex02/20_newsgroups/"

```

```

#folders in the path
folders = os.listdir(path)
#change the current working directory to path
os.chdir(path)

final_list = []
#going through each folders and files inside of them.
for i in range(len(folders)):
    line_inner = []
    filenames = []
    inner_path = (path+str(folders[i]))
    os.chdir(inner_path)
    filenames += [f for f in os.listdir(inner_path) if os.path.isfile(f)]
    #determining the share for each process
    share = round(len(filenames)/size)

    #only slaves
    if rank != 0:

        #receiving names of the files
        filens = comm.recv(source=0, tag = 0)
        #assigning files to ditexterent workers except worker 0

        #cleaning and tokenization
        tokenized = clean_tokenized(filens)

        #sending to master node
        comm.send(tokenized, dest=0, tag=1)

    #master process
    else:

        #distributing the data
        for i in range(1, size):

            comm.send(filenames[(share*i):(share*(1+i))], dest = i, tag = 1)

        #master node's share of file
        filenames = filenames[(share*rank):(share*(1+rank))]
        #cleaning and tokenization
        tokenized = clean_tokenized(filenames)
        #saving the file
        #with open("/home/mansoor/Desktop/DDA/ex02/saved"+f'/tokenized_{rank}.pkl', 'wb') as fp:
        #    pickle.dump(tokenized, fp)
        #appending
        line_inner.append(tokenized)

        #receiving results
        for i in range(1, size):

            tknzd = comm.recv(source = i, tag = 1)

            line_inner.append(tknzd)

        #appending final result
        final_list.append(line_inner)

```

```
t1 = MPI.Wtime() - t0
print("Time: ", t1)
#print(len(final_list))
#save_csv(final_list, rank=10)
with open("/home/mansoor/Desktop/DDA/ex02/saved"+f'/final_list_{size}.pkl', 'wb') as fp:
    pickle.dump(final_list, fp)
return final_list

main_func = main_function()
```

Overwriting DDA02_ex1.py

In [8]: `!mpiexec -n 4 python DDA02_ex1.py`

```
Time: 25.713739156723022
Time: 25.72397494316101
Time: 25.755807876586914
Time: 25.765300989151
```

Exercise02

Report 2

Flowchart of the processes

```
-Initialization
Loading saved data
```

Table of Times for 2

P=2	P=4	P=6
2.9	2.3	6.02

- The result of parallelization visible in this exercise
- As the number of workers increases, the computation time decreases but only until 4.
- our best result is when we use 4 workers.

Strategy

- We loop over the files of each folder.
- we share the files in each folder to different workers.
- each worker will perform tf computation on the subset of data it has and send the result back to rank0.
- rank0 is our master process and it also work as worker, but at the end it receives all the data and add them in one big list.

```
[{'path': 0.0072992700729927005,
  'cantaloupe': 0.0072992700729927005,
  'news': 0.014598540145985401,
  'harvard': 0.0072992700729927005,
  'near': 0.0072992700729927005,
  'howland': 0.0072992700729927005,
  'reston': 0.0072992700729927005,
  'james': 0.014598540145985401,
  'felder': 0.014598540145985401,
  'newsgroups': 0.0072992700729927005,
  'atheism': 0.014598540145985401,
  'subject': 0.0072992700729927005,
  'help': 0.021897810218978103,
  'god': 0.021897810218978103,
  'court': 0.0072992700729927005,
  'date': 0.0072992700729927005,
  'apr': 0.014598540145985401,
  'gmt': 0.0072992700729927005,
  'organization': 0.0072992700729927005,
```

result of exercise 2

In [176...

```

%%writefile ex02_tf.py

from mpi4py import MPI
import numpy as np
import os
import pickle

#using saved data
with open ('saved/final_list_1.ob','rb') as fp:
    data = pickle.load(fp)

#initialization
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()
#name = MPI.Get_processor_name()

#batch = rank(len(data-1)/size)

#tf calculator function
def calculate_tf(tokenized_data):
    tf_list = []
    for document in tokenized_data:
        sentence_dict = dict()
        for word in document:
            sentence_dict[word] = sentence_dict.get(word,0)+1
        len_docu = len(document)

        for word in sentence_dict:
            sentence_dict[word] = sentence_dict[word]/len_docu
        tf_list.append(sentence_dict)
    return tf_list

t0 = MPI.Wtime()
def main_function():
    """
    This function is the main function to run the parallel program on.
    """
    #time

    final_list = []
    #going through each folders and files inside of them.
    for i in range(len(data)-1):
        line_inner = []

        #determining the share for each process
        share = round(len(data[i][0])/size)
        #the data is selected
        selected = data[i][0]

        #only slaves
        if rank != 0:

            #receiving names of the files
            filens = comm.recv(source=0, tag = 0)

```

```

    #assigning files to ditexterent workers except worker 0

    #calculating tf
    tf_res = calculate_tf(filens)

    #sending to master node
    comm.send(tf_res, dest=0, tag=1)

#master process
else:

    #distributing the data
    for i in range(1, size):

        comm.send(selected[(share*i):(share*(1+i))], dest = i, tag=1)

    #master node's share of file
    filens = selected[(share*rank):(share*(1+rank))]
    #tf
    tf_res = calculate_tf(filens)
    line_inner.append(tf_res)

    #receiving results
    for i in range(1, size):

        tff = comm.recv(source = i, tag = 1)
        line_inner.append(tff)

    #appending final result
    final_list.append(line_inner)

t1 = MPI.Wtime() - t0
print("Time: ", t1)

with open("/home/mansoor/Desktop/DDA/ex02/saved"+f'/tf_res_{size}.ob',
          pickle.dump(final_list, fp)
return final_list

main_func = main_function()

```

Overwriting ex02_tf.py

In [174...

```
!mpiexec -n 4 python ex02_tf.py
```

Time: 2.2453770637512207

Time: 2.2325971126556396

Time: 2.3546640872955322

Time: 2.2729651927948

In [121...

```

#just loading the data of tf
import pickle

#using saved data
with open ('saved/final_list_1.ob','rb') as fp:
    data = pickle.load(fp)

```


Exercise 3

Report 3

Table of Times for 3

P=2	P=4	P=6
7.9796	5.9459	6.0164

- The result of parallelization visible in this exercise
 - As the number of workers increases, the computation time decreases.
 - our best result is when we use 4 workers.
-

Strategy

- first we use the tokenized data from previous exercise.
- we append all of the documents in one list.
- word_frequency_in_doc() function goes through the data find the occurrence and then ratio of word in the each batch.
- the result sent to the master process and it does the final frequency.
- save the data at the end

```
{'affirmation': 7.7448717855918305,
'allah': 5.888970290173582,
'alternate': 5.668168390698289,
'andrew': 1.58557334182599,
'another': 1.9615282663449052,
'apr': -0.2049835364304671,
'arizona': 4.073086895287181,
'article': 0.49683519017798505,
'ask': 2.6628640892822806,
'asked': 3.1646900611867914,
'atheism': 1.8497920501626317,
'back': 1.856417812569192,
'cantaloupe': -0.2876820724517809,
'case': 2.075684688582293,
'center': 2.590313975395854,
'cleveland': 3.5812408758914436,
'court': 4.089015924808619,
'daffy': 6.194829671960989,
'date': -0.2876820724517809,
```

result of exercise 3

In [161]...

```

%%writefile ex03_idf_0.py

#loading libraries
from mpi4py import MPI
import numpy as np
import pandas as pd
import pickle
import os
from collections import defaultdict, Counter

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

#using saved data
with open ('saved/final_list_1.ob','rb') as fp:
    data = pickle.load(fp)
    total_docs = []
    for i in range(len(data)-1):
        for j in range(len(data[i][0])):
            total_docs.append(data[i][0][j])

def word_frequency_in_doc(data):
    word_folder = dict()
    for i in range(len(data)):
        # Get unique words per document
        words = np.unique(data[i])
        # Counting the times a word has been mentioned in a document
        for word in words:
            if word in word_folder:
                word_folder[word] += 1
            else:
                word_folder[word] = 1
    batch = {k: (len(data) / v) for k, v in word_folder.items()}
    return batch

p_ranks = round(len(total_docs)/(size-1))

total = dict()

#master process
if rank == 0:
    t0 = MPI.Wtime()
    total_docs2 = total_docs[0:p_ranks]
    for i in range(1, size):
        total_docs1 = total_docs[(i*p_ranks):(p_ranks*(i+1))]
        comm.send(total_docs1, dest=i)
    output1 = word_frequency_in_doc(total_docs2)

    global_dict = None
    for i in range(1, size):
        idf = comm.recv()
        output1 = (Counter(output1) + Counter(idf))
    total = {k: np.log(v / (size)) for k, v in output1.items()}

    print('Time:', MPI.Wtime() - t0)
    with open ("/home/mansoor/Desktop/DDA/ex02/saved"+f"/idf_res_{size}.ob"
              pickle.dump(total, fp)

```

```
#slaves
else:
    data = comm.recv()
    output = word_frequency_in_doc(data)
    comm.send(output, dest=0)
```

Overwriting ex03_idf_0.py

In [162... `!mpiexec -n 4 python ex03_idf_0.py`

Time: 5.945959806442261

In [139... `#using saved data`
`with open ('saved/idf_res_4.ob', 'rb') as fp:`
`idf_ = pickle.load(fp)`

Exercise 4

Report 4

Table of Times for 4

P=2	P=4	P=6
15.98	12.679	17.332

- The result of parallelization visible in this exercise
 - As the number of workers increases, the computation time decreases.
 - our best result is when we use 4 workers.
-

Strategy

- first we use the tokenized data from previous exercise.
- we append all of the documents in one list.
- the master process sends batches of data to the remaining workers.
- each worker will accomplish 2 tasks: first get the TF for each document and secondly it will get the IDF for all the documents that each has received from the root.
- Master will merge the data from all the processes, measure the final IDF and multiply with each token TF.
- save the data at the end.

```
{'path': 0.0,  
 'cantaloupe': 0.0,  
 'news': 0.0005527434658239272,  
 'harvard': 0.0036318256526451614,  
 'near': 0.006778819166329211,  
 'howland': 0.0018767247696019483,  
 'reston': 0.0018770432215244105,  
 'james': 0.010002905866735473,  
 'felder': 0.1878706090841518,  
 'newsgroups': 0.0,  
 'atheism': 0.012915251496159595,  
 'subject': 0.0,  
 'help': 0.02934840813854678,  
 'god': 0.08426704390159284,  
 'court': 0.013222652559699092,  
 'date': 0.0,  
 'apr': 0.000255242395127512,  
 'gmt': 0.0004976714001729526,  
 'organization': 0.0002552293516974751,
```

result of exercise 4

In [187...

```

%%writefile ex04_tf_idf.py

#loading libraries
from mpi4py import MPI
import numpy as np
import pandas as pd
import pickle
import os
from collections import defaultdict, Counter

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

#using saved data
with open ('saved/final_list_1.ob', 'rb') as fp:
    data = pickle.load(fp)
    total_docs = []
    for i in range(len(data)-1):
        for j in range(len(data[i][0])):
            total_docs.append(data[i][0][j])

def calculate_tfidf(data):
    chunk = []
    word_folder = dict()
    for i in range(len(data)):
        word = defaultdict(int)
        unique_tf = np.unique(data[i])
        counter = len(np.unique(unique_tf))
        for i in data[i]:
            word[i] +=1
        dictionary = {k: v / counter for k, v in word.items()}
        chunk.append(dictionary)
        for unique in unique_tf:
            if unique in word_folder:
                word_folder[unique] += 1
            else:
                word_folder[unique] = 1
    batch = {k: (len(data) / v) for k, v in word_folder.items()}

    return chunk, batch

p_ranks = round(len(total_docs)/(size-1))
final_idf = dict()
final = []
dictcalculate_tfidf = dict()
#master process
if rank == 0:
    t0 = MPI.Wtime()
    total_docs2 = total_docs[0:p_ranks]
    for i in range(1, size):
        total_docs1 = total_docs[(i*p_ranks):(p_ranks*(i+1))]
        comm.send(total_docs1, dest=i)
    TF1, IDF1 = calculate_tfidf(total_docs2)

```

```

global_dict = None
for i in range(1, size):
    chunka = []
    TF2, IDF2 = comm.recv()
    TF1 = TF1 + TF2
    IDF1 = (Counter(IDF2) + Counter(IDF1))
final = TF1 + final
final_idf = {k: np.log(v / (size-1)) for k, v in IDF1.items()}
for i in range(len(final)):
    n = final[i]
    for key, value in n.items():
        if key in final_idf:
            dictcalculate_tfidf[key] = (n[key]) * (final_idf[key])
        else:
            None

print('Time:', MPI.Wtime() - t0)

else:
    data = comm.recv()
    TF, IDF = calculate_tfidf(data)
    output = (TF, IDF)
    comm.send(output, dest=0)

with open("/home/mansoor/Desktop/DDA/ex02/saved"+f'/tfidf_res_{size}.ob',
          pickle.dump(dictcalculate_tfidf, fp)

```

Overwriting ex04_tf_idf.py

In [188... `!mpirun -n 4 python ex04_tf_idf.py`

Time: 12.67906403541565

In [197... `#using saved data`
`with open ('saved/tfidf_res_4.ob', 'rb') as fp:`
`tfidf_ = pickle.load(fp)`

In []: