

Mansoor Nabawi, 309498

References:

- <https://www.askpython.com/python-modules/pytorch-custom-datasets>
- <https://debuggercafe.com/custom-dataset-and-dataloader-in-pytorch/>

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch.utils.data as data

import torchvision.transforms as transforms
import torchvision.datasets as datasets

from sklearn import decomposition
from sklearn import manifold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
from tqdm.notebook import tqdm, trange
import matplotlib.pyplot as plt
import numpy as np

import copy
import random
import time

from torch.utils.tensorboard import SummaryWriter
%load_ext tensorboard

SEED = 1234

random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

loading and normalizing data

```
ROOT = '.data'

train_data = datasets.MNIST(root=ROOT,
                             train=True,
                             download=True)

mean = train_data.data.float().mean() / 255
std = train_data.data.float().std() / 255

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to .data/MNIST/raw/train-images-idx3-ubyte.gz
9913344/? [00:00<00:00, 53830832.94it/s]

Extracting .data/MNIST/raw/train-images-idx3-ubyte.gz to .data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to .data/MNIST/raw/train-labels-idx1-ubyte.gz
29696/? [00:00<00:00, 1014746.56it/s]

Extracting .data/MNIST/raw/train-labels-idx1-ubyte.gz to .data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to .data/MNIST/raw/t10k-images-idx3-ubyte.gz
1649664/? [00:00<00:00, 34941356.88it/s]

Extracting .data/MNIST/raw/t10k-images-idx3-ubyte.gz to .data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to .data/MNIST/raw/t10k-labels-idx1-ubyte.gz
5120/? [00:00<00:00, 153759.61it/s]

Extracting .data/MNIST/raw/t10k-labels-idx1-ubyte.gz to .data/MNIST/raw
```

defining some transformations for test and train

```
train_transforms = transforms.Compose([
    transforms.RandomRotation(5, fill=(0,)),
    transforms.RandomCrop(28),
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])

test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean=[mean], std=[std])
])

train_data = datasets.MNIST(root=ROOT,
                             train=True,
                             download=True,
                             transform=train_transforms)

test_data = datasets.MNIST(root=ROOT,
                             train=False,
                             download=True,
                             transform=test_transforms)
```

writing the custom dataset

- I am going to get the dataset in my custom class.
- we make random indexes of data but in a smaller size.
- we choose the random indexes of data and put it in another variables.
- Then i will concat/stack based on the number of k we have.

- 1later this class will be used by train iterator.

```
import numpy as np
from torch.utils.data import Dataset

class custom_data(Dataset):
    def __init__(self, data, k):
        self.data = data
        self.lendata = len(self.data)
        self.k = k
        self.get_index()

    def __len__(self):
        return len(self.data)//self.k

    def __getitem__(self, idx):
        indx = self.indx_[idx]
        x = []
        y = []
        for i in indx:
            x.append(self.data.__getitem__(i)[0])
            y.append(self.data.__getitem__(i)[1])

        data_, label = torch.stack(x), np.sum(y).astype(dtype=np.float32)
        return data_, label

    def get_index(self):
        self.indx_ = np.random.randint(0,self.lendata, size=(self.lendata//self.k, self.k))
```

```
np.random.randint(0,len(train_data), size=(len(train_data)//3, 3)).shape

(20000, 3)
```

```
train_cdata = custom_data(train_data,k = 3)
test_cdata = custom_data(test_data,k = 3)
```

```
len(train_cdata)

20000
```

```
BATCH_SIZE = 64

train_iterator = data.DataLoader(train_cdata,
                                shuffle=True,
                                num_workers=2,
                                batch_size=BATCH_SIZE)

test_iterator = data.DataLoader(test_cdata,
                                num_workers=2,
                                batch_size=BATCH_SIZE,
                                shuffle = False)
```

```
example = enumerate(train_iterator)
idx, (dt, lb) = next(example)
```

N,K,C,W,H

```
dt.shape

torch.Size([64, 3, 1, 28, 28])
```

```
lb.shape

torch.Size([64])
```

we make few changes to our NN.

we add another layer for the final prediction, as we have ten numbers but at the end we will have sum of the three numbers and it is going to be only one output.

```
class LeNet(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(in_channels=1,
                                out_channels=6,
                                kernel_size=5,
                                padding=2)

        self.conv2 = nn.Conv2d(in_channels=6,
                                out_channels=16,
                                kernel_size=5)

        self.fc_1 = nn.Linear(16 * 5 * 5, 120)
        self.fc_2 = nn.Linear(120, 84)
        self.fc_3 = nn.Linear(84, 10)
        self.fc_4 = nn.Linear(10,1)

    def forward(self, x):
        N,K,C,W,H = x.size()
        x = x.view(N*K,C,W,H)

        x = self.conv1(x)
        x = F.max_pool2d(F.relu(x), kernel_size=2)

        x = self.conv2(x)
        x = F.relu(x)

        x = F.max_pool2d(x, kernel_size=2)
        x = x.view(-1,16*5*5 )

        x = self.fc_1(x)
        x = F.relu(x)
```

```
        x = self.fc_2(x)
        x = F.relu(x)

        x = self.fc_3(x)
        x = F.relu(x)

        x = self.fc_4(x)
        x = F.relu(x)
        x = x.view(N,K)
        x = torch.sum(x, dim=1)

    return x
```

```
print(LeNet())
```

```
LeNet(
  (conv1): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc_1): Linear(in_features=400, out_features=120, bias=True)
  (fc_2): Linear(in_features=120, out_features=84, bias=True)
  (fc_3): Linear(in_features=84, out_features=10, bias=True)
  (fc_4): Linear(in_features=10, out_features=1, bias=True)
)
```

```
def train(model, iterator, optimizer, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.train()

    for (x, y) in tqdm(iterator, desc="Training", leave=False):

        x = x.to(device)
        y = y.to(device)

        optimizer.zero_grad()

        y_pred = model(x)

        loss = criterion(y_pred, y)

        loss.backward()

        optimizer.step()

        epoch_loss += loss.item()

    return epoch_loss / len(iterator)
```

```
def test(model, iterator, criterion, device):

    epoch_loss = 0
    epoch_acc = 0

    model.eval()

    with torch.no_grad():

        for (x, y) in tqdm(iterator, desc="Testing", leave=False):

            x = x.to(device)
            y = y.to(device)

            y_pred = model(x)

            loss = criterion(y_pred, y)

            epoch_loss += loss.item()

    return epoch_loss / len(iterator)
```

```
def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs
```

```
logs = 'runs/mnist_custom'
tb = SummaryWriter(logs)
```

```
selected_optimizer='Adam'
for learning_rate in [0.001]:

    model = LeNet()

    optimizer = optim.Adam(model.parameters())

    criterion = nn.MSELoss()

    #Checking if we can use GPU

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    model = model.to(device)
    criterion = criterion.to(device)

    avail_optimizers = {'Adam':torch.optim.Adam(model.parameters(), lr=learning_rate, betas=(0.9, 0.999), eps=1e-08, weight_decay=0, amsgrad=False),
                        'RMS': torch.optim.RMSprop(model.parameters(), lr=learning_rate, alpha=0.99, eps=1e-08, weight_decay=0, momentum=0, centered=False),
                        'SGD': torch.optim.SGD(model.parameters(), lr=learning_rate, momentum=0, dampening=0, weight_decay=0, nesterov=False)}

    optimizer = avail_optimizers[selected_optimizer]

    epochs = 20
```

```
print(f"selected optimizer and learning rate: {optimizer}")
for n_iter in tqdm(range(epochs)):
    print(f"Epoch {n_iter+1}/{epochs}")

    start_time = time.monotonic()

    train_loss = train(model, train_iterator, optimizer, criterion, device)
    train_iterator.dataset.get_index()
    test_loss = test(model, test_iterator, criterion, device)
    test_iterator.dataset.get_index()

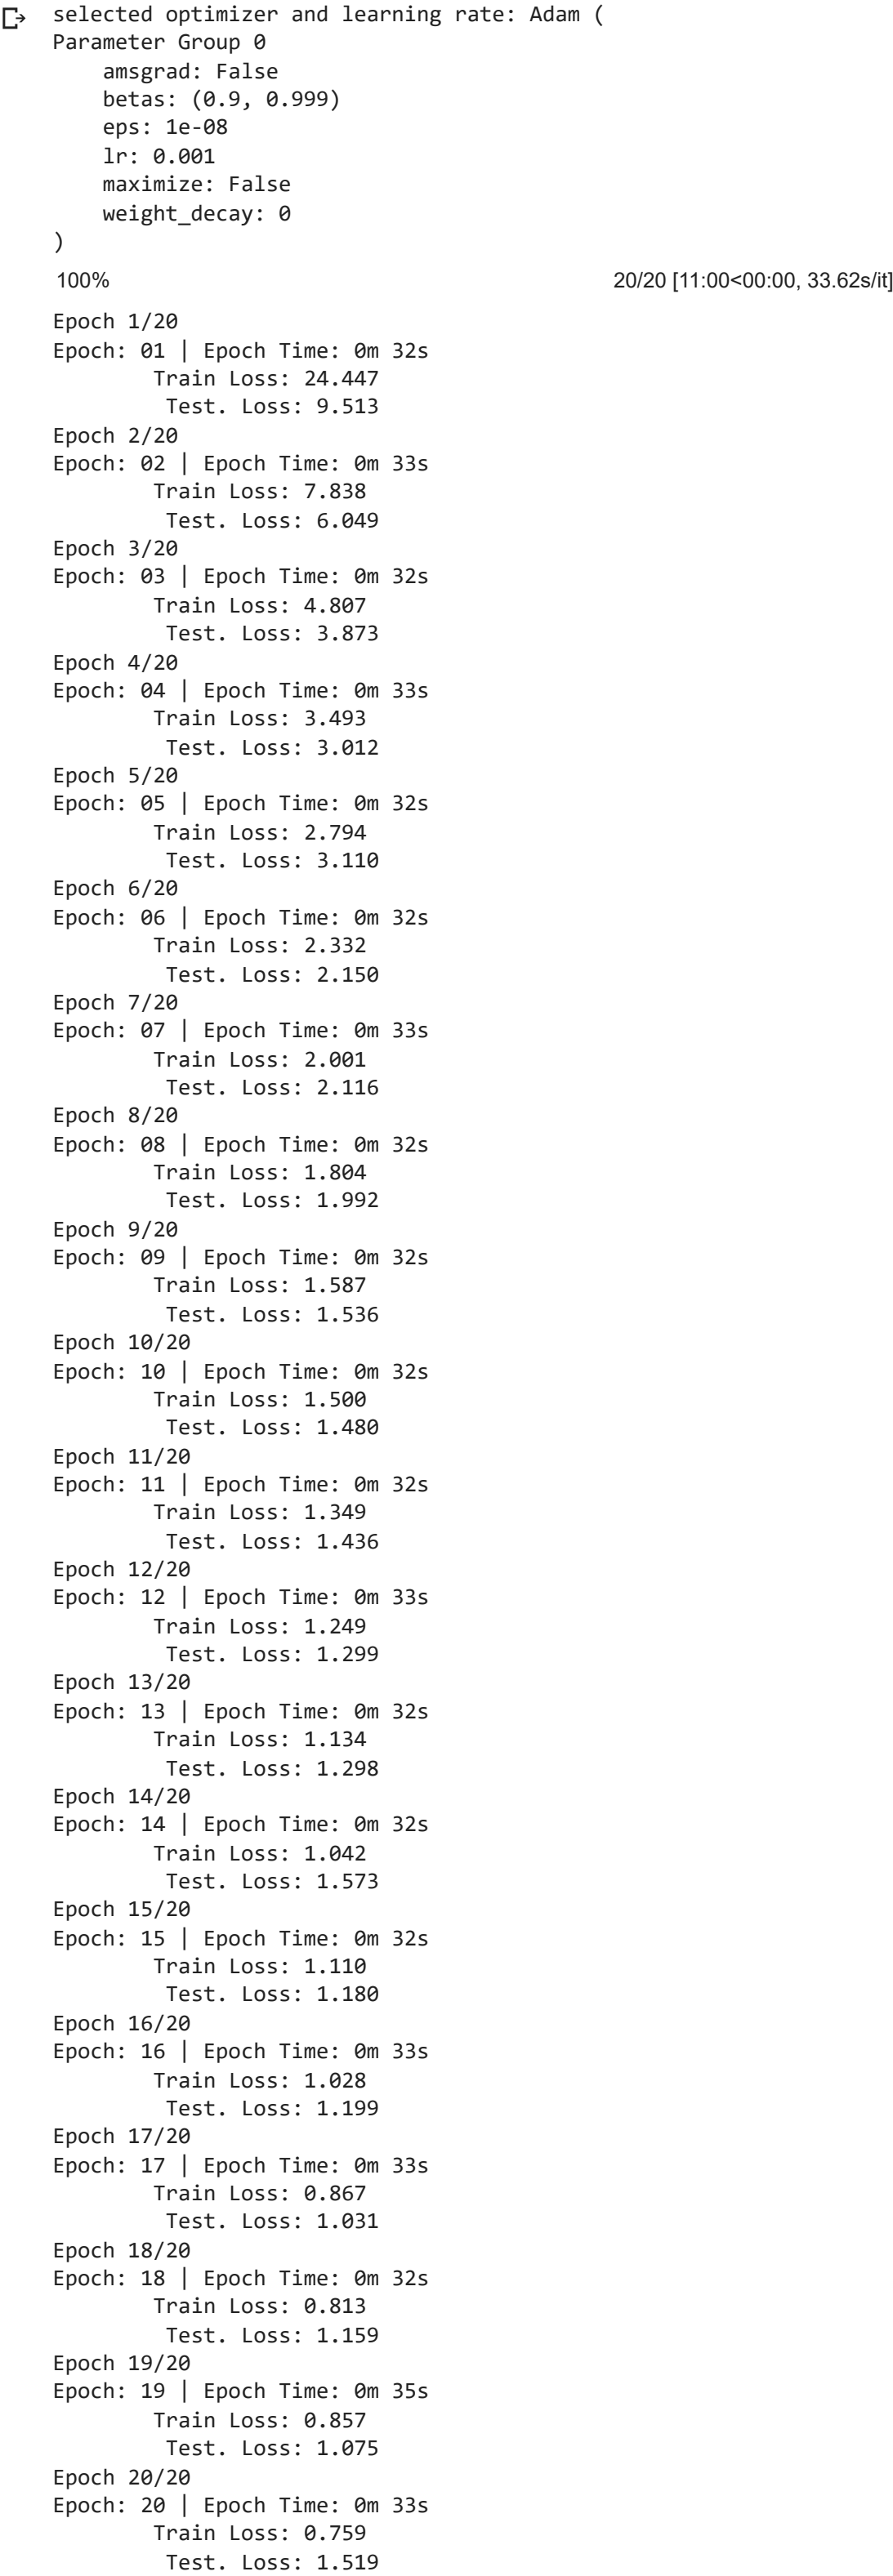
    end_time = time.monotonic()

    epoch_mins, epoch_secs = epoch_time(start_time, end_time)

    print(f'Epoch: {n_iter+1:02} | Epoch Time: {epoch_mins}m {epoch_secs}s')
    print(f'\tTrain Loss: {train_loss:.3f} ')
    print(f'\tTest. Loss: {test_loss:.3f} ')

    if tb is not None:
        tb.add_scalars(f'Loss_{learning_rate}', {"Train":train_loss,
                                                "Test":test_loss}, n_iter)

        tb.add_scalars(f'Train_loss', {f"{learning_rate}":train_loss}, n_iter)
        tb.add_scalars(f'Test_loss', {f"{learning_rate}":test_loss}, n_iter)
```



As it is visible our model is performing very well and the loss is decreasing

```
%tensorboard --logdir runs/mnist_custom
```

